

# CREATING & USING A LIVE, LOCALIZED EXTRACT OF OPENSTREETMAP

# OVERVIEW

- Introduction
- OSM Data
- Import Scenarios
- Keeping it Up to Date
- Using it
- Questions

# STEAL THIS PRESENTATION!



# OPENSTREETMAP DATA

# WHAT IS OSM?

*"The Wikipedia of Maps"*

OSM is a **dataset**, which is...

- global
- open
- used to map nearly anything
- one of my favorite things


# IS ANYONE ACTUALLY USING IT?

- Amazon Logistics
- Tesla Smart Summon feature
- Esri (basemaps, feature services)
- Pokemon Go
- Red Cross
- Kendall County!

# OPENSTREETMAP DATA: *IT'S WEIRD!*

The OSM dataset has:

- NO layers
- NO schema
- NO polygons



To be more precise, OSM is an **object-oriented data model**.  
Every element in the dataset can have an arbitrary number of key-value  
pairs

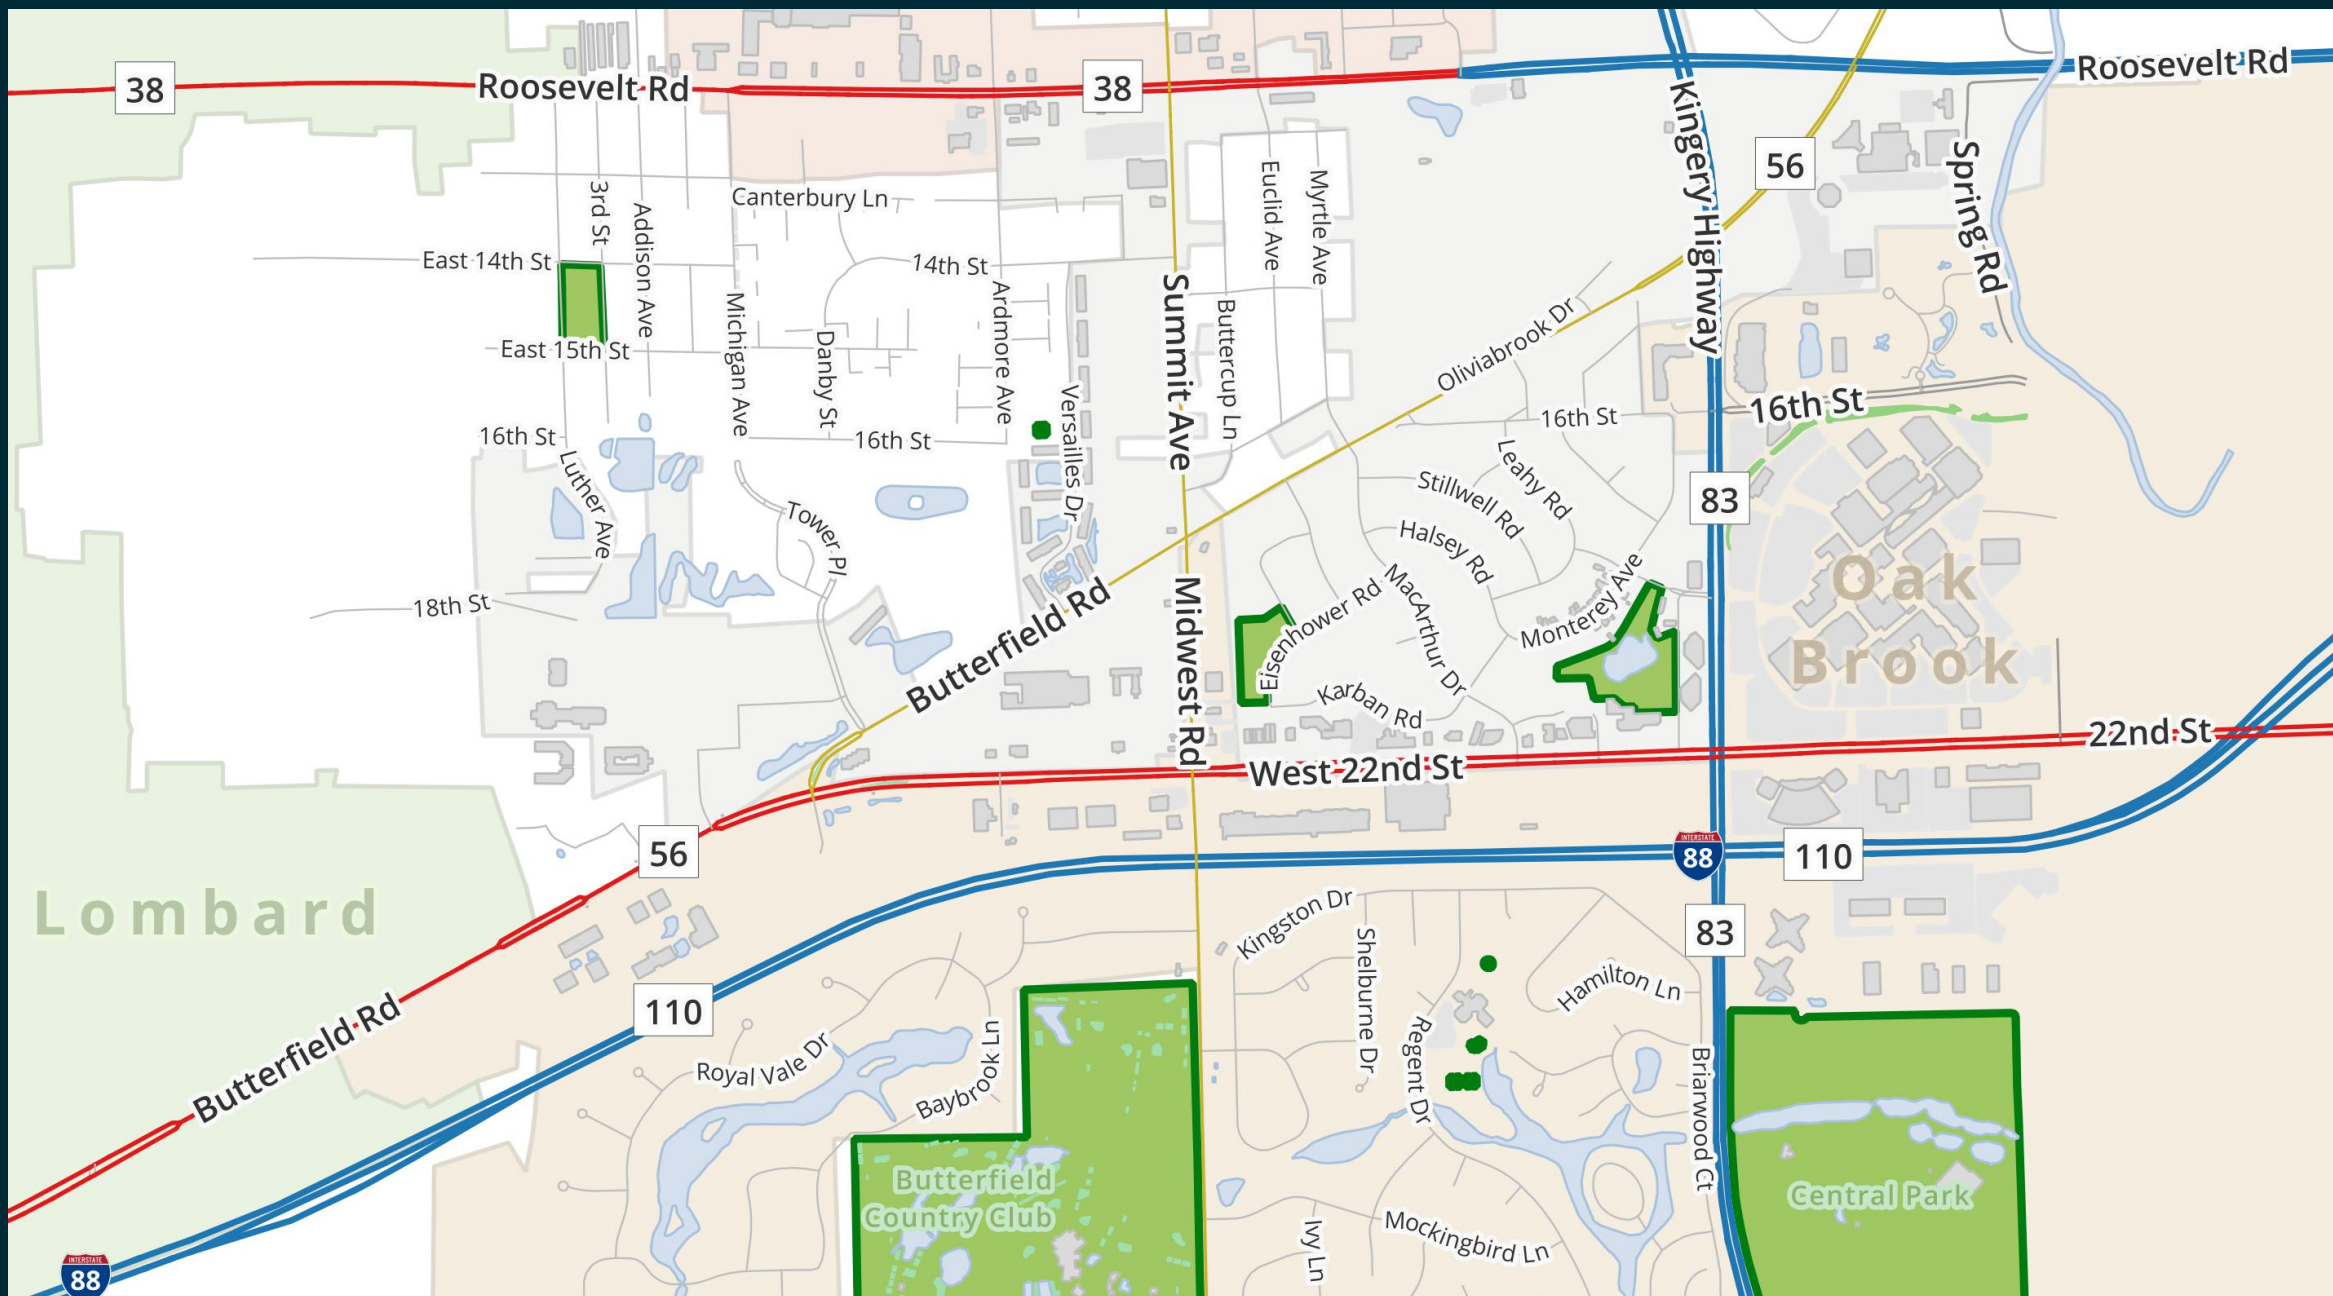
In OSM, we call those "tags".

```
1 {  
2   "elements": [  
3     {  
4       "type": "node",  
5       "id": 1,  
6       "tags": {  
7         "amenity": "bench"  
8       }  
9     },  
10    {  
11      "type": "way",  
12      "id": 2,  
13      "tags": {  
14        "natural": "water",  
15        "water": "lake",
```









# GETTING THE DATA

# SOURCES

[https://wiki.openstreetmap.org/wiki/Downloading\\_data](https://wiki.openstreetmap.org/wiki/Downloading_data)

For very small areas, it may be possible to query the **Overpass API** directly.

State-based extract from GeoFabrik:

<https://download.geofabrik.de/north-america/us/illinois.html>

# OSMIUM

<https://osmcode.org/osmium-tool/>

A "swiss army knife" for OSM data.

- **Extracting geographic subset**
- Getting metadata about objects
- Extract by attribute
- Generate / apply change files
- and more!

# AUDIENCE PARTICIPATION!

<http://geojson.io/#map=6.2/39.976/-89.164>



# IMPORTING



# imposm

## Pros

- ✓ Easy to use and configure
- ✓ Built in tag "cleaning"
- ✓ Table generalization
- ✓ Area of interest filtering for import *and* updating

## Cons

- ✗ Relations are harder to work with, and require own tables
- ✗ Lacks advanced geometry processing or configuration
- ✗ Area of interest does not apply to relations
- ✗ Development has plateaued

# Mapping OSM data to tables is done using a YAML config file.

```
1 tables:
2   roads:
3     columns:
4       - name: osm_id
5         type: id
6       - name: the_geom
7         type: geometry
8       - name: type
9         type: mapping_value
10      - key: name
11        name: name
12        type: string
13    filters:
14      reject:
15        area: ["yes"]
16    mapping:
```

# IMPORTING: SINGLE TABLE

```
1 imposm import \  
2   -config /app/imposm-scenarios/highways-config.json \  
3   -read /app/data/city-extract.osm.pbf \  
4   -deployproduction -optimize -write -overwritecache
```

```
1 {  
2   "cachedir": "/app/data/cache/highways",  
3   "mapping": "/app/imposm-scenarios/highways-mapping.yml",  
4   "connection": "postgis://gis:gis@database:5432/ilgisa2023?prefix=state_highw  
5 }
```

# ADDITIONAL SCENARIOS

## MULTIPLE TABLES

```
imposm import \  
  -config /app/imposm-scenarios/city-parks-config.json \  
  -read /app/data/city-extract.osm.pbf \  
  -deployproduction -optimize -write -overwritecache
```

## IMPORTING *EVERYTHING*

And prepping for future updates!

```
1 imposm import \  
2   -config /app/imposm-scenarios/city-all-config.json \  
3   -read /app/data/city-extract.osm.pbf \  
4   -deployproduction -optimize -write -overwritecache \  
5   -diff
```

# osm2pgsql

## Pros

✓ Well established,  
continued development

✓ Extremely  
configurable

✓ Allows mid-stream  
geometry operations

✓ Exports to many  
coordinate systems

## Cons

✗ Updating requires a replication url,  
cannot easily be limited to an area of  
interest

✗ Custom configuration can be harder  
to adjust / understand

✗ Built-in generalization is still only  
experimental

## Mapping to tables is done with a Lua file.

```
1 local nodes = osm2pgsql.define_table({
2     name = 'osm2pgsql_nodes',
3     ids = { type = 'any', id_column = 'osm_id' },
4     columns = {
5         { column = 'id', sql_type = 'serial', create_only = true },
6         { column = 'tags', type = 'hstore' },
7         { column = 'the_geom', type = 'point', not_null = true }
8     }
9 })
10
11 function osm2pgsql.process_node(object)
12     nodes:insert({
13         tags = object.tags,
14         the_geom = object:as_point()
15     })
16 end
```

# IMPORTING

## POINTS OF INTEREST

```
1 osm2pgsql \  
2   -j /app/data/city-extract.osm.pbf \  
3   -d postgres://gis:gis@database:5432/ilgisa2023 \  
4   -O flex -S /app/osm2pgsql-scenarios/pois.lua
```

## EVERYTHING

```
osm2pgsql \  
  -j /app/data/city-extract.osm.pbf \  
  -d postgres://gis:gis@database:5432/ilgisa2023 \  
  -O flex -S /app/osm2pgsql-scenarios/city-all.lua
```

# UPDATING

```
imposm run -config /app/imposm-scenarios/city-all-config.json
```

Updating with `osm2pgsql` will only pull apply updates from the original download, i.e., all changes in the GeoFabrik IL extract, or all minutely changes for the entire planet.



# USING IT

# THOSE "EVERYTHING" TABLES

Rather than filter the data on import, it is possible to import all elements.

The tags can all be put into a single `jsonb` column. Attributes stored in the following format:

```
{"key_1":"value_1", "key_2":"value_2", ... "key_n":"value_n"}
```

Using the `column -> 'key'` syntax returns the value for the specified key.

```
1 SELECT
2     osm_id,
3     the_geom,
4     tags -> 'cuisine' as cuisine
```

## WHY WOULD I WANT THAT?

- Schema remains as flexible as OSM itself
- No re-importing / re-indexing
- As use cases change, only the *queries* need to be modified, not the data
- Query output can be used identically to any SQL table
- QGIS has built-in support to parse *and edit* `jsonb` fields

# SOME QUERIES

## TRAFFIC ROADS

```
SELECT
  osm_id,
  the_geom,
  tags -> 'name' name,
  tags -> 'highway' class
FROM osm2pgsql_ways
WHERE tags -> 'highway' IN (
  'primary',
  'secondary',
  'tertiary',
  'motorway',
  'trunk',
  'unclassified',
  'residential',
  'service'
)
```

## NATURAL AREAS / LAND COVER

```
SELECT
    osm_id,
    the_geom,
    COALESCE(tags -> 'natural', tags->'landuse') type,
    tags
FROM kendall_areas
WHERE COALESCE(tags -> 'natural', tags->'landuse') IS NOT NULL
```

# WHAT TO DO

Honestly, once you write the query, you can use the results for **anything** that you would a normal table, including:

- Publishing a feature service to the web
- Generating vector tiles
- Using as input in geoprocessing
- Create an ArcGIS Locator dataset
- Create a routable graph network

# QUESTIONS?